

TOPIC : **"Explainable AI using class activation mappings in Pytorch"**

GROUP NUMBER : 7

TEAM MEMBERS : **Kashif Riyaz, Nitesh Morem**

UNDER THE GUIDANCE : **Prof. Dr.-Ing. Christian Bergler**

We Used the Jacob gills github Respository while working on the code.

The Repository Link is "<https://github.com/jacobgil/pytorch-grad-cam>"

0 Installing necessary Modules

In this sections all the necessary modules are instralled which include:

- Pytorch
- torchvision
- numpy
- opencv
- matplotlib
- grad-cam

Note: The Grad-CAM modules have been imported from the jacob gill's github repository: reference link: <https://github.com/jacobgil/pytorch-grad-cam>

```
In [91]: !pip install --upgrade pip
!pip install torch torchvision numpy opencv-python requests pillow grad-cam matp
```

```
Requirement already satisfied: pip in c:\users\kashi\virtualenv\ai_projects\lib\s
ite-packages (24.0)
```

```
Collecting pip
```

```
Using cached pip-24.1.1-py3-none-any.whl.metadata (3.6 kB)
```

```
Using cached pip-24.1.1-py3-none-any.whl (1.8 MB)
```

```
ERROR: To modify pip, please run the following command:
```

```
C:\Users\kashi\virtualenv\AI_Projects\Scripts\python.exe -m pip install --upgrade
pip
```

```
[notice] A new release of pip is available: 24.0 -> 24.1.1
```

```
[notice] To update, run: python.exe -m pip install --upgrade pip
```

Requirement already satisfied: torch in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (2.2.2)

Requirement already satisfied: torchvision in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (0.17.2)

Requirement already satisfied: numpy in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (1.26.4)

Requirement already satisfied: opencv-python in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (4.9.0.80)

Requirement already satisfied: requests in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (2.31.0)

Requirement already satisfied: pillow in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (10.3.0)

Requirement already satisfied: grad-cam in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (1.5.0)

Requirement already satisfied: matplotlib in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (3.8.4)

Requirement already satisfied: filelock in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from torch) (3.13.4)

Requirement already satisfied: typing-extensions>=4.8.0 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from torch) (4.11.0)

Requirement already satisfied: sympy in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from torch) (1.12)

Requirement already satisfied: networkx in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from torch) (3.3)

Requirement already satisfied: jinja2 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from torch) (3.1.3)

Requirement already satisfied: fsspec in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from torch) (2024.3.1)

Requirement already satisfied: charset-normalizer<4,>=2 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from requests) (3.3.2)

Requirement already satisfied: idna<4,>=2.5 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from requests) (3.7)

Requirement already satisfied: urllib3<3,>=1.21.1 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from requests) (2.2.1)

Requirement already satisfied: certifi>=2017.4.17 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from requests) (2024.2.2)

Requirement already satisfied: ttach in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from grad-cam) (0.0.3)

Requirement already satisfied: tqdm in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from grad-cam) (4.66.2)

Requirement already satisfied: scikit-learn in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from grad-cam) (1.4.2)

Requirement already satisfied: contourpy>=1.0.1 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from matplotlib) (1.2.1)

Requirement already satisfied: cycler>=0.10 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from matplotlib) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from matplotlib) (4.51.0)

Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from matplotlib) (1.4.5)

Requirement already satisfied: packaging>=20.0 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from matplotlib) (24.0)

Requirement already satisfied: pyparsing>=2.3.1 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from matplotlib) (3.1.2)

Requirement already satisfied: python-dateutil>=2.7 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from matplotlib) (2.9.0.post0)

Requirement already satisfied: six>=1.5 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from python-dateutil>=2.7->matplotlib) (1.16.0)

Requirement already satisfied: MarkupSafe>=2.0 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from jinja2->torch) (2.1.5)

Requirement already satisfied: scipy>=1.6.0 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from scikit-learn->grad-cam) (1.13.0)
 Requirement already satisfied: joblib>=1.2.0 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from scikit-learn->grad-cam) (1.4.0)
 Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from scikit-learn->grad-cam) (3.4.0)
 Requirement already satisfied: mpmath>=0.19 in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from sympy->torch) (1.3.0)
 Requirement already satisfied: colorama in c:\users\kashi\virtualenv\ai_projects\lib\site-packages (from tqdm->grad-cam) (0.4.6)

[notice] A new release of pip is available: 24.0 -> 24.1.1

[notice] To update, run: python.exe -m pip install --upgrade pip

1. Importing necessary modules

In this sections all the modules are imported

```
In [92]: import torchvision
import warnings
import torch
warnings.filterwarnings('ignore')
from torchvision import models
import numpy as np
import cv2
import requests
from PIL import Image

import matplotlib.pyplot as plt
```

Note: if there is an error regarding any module please restart the kernal and run the whole code again

1.1 Importing modules from pytorch_grad_cam

```
In [93]: from pytorch_grad_cam.utils.model_targets import ClassifierOutputTarget
from pytorch_grad_cam.utils.image import show_cam_on_image, \
    deprocess_image, \
    preprocess_image
```

1.1.1 Importing modules for visualisation

```
In [94]: from pytorch_grad_cam.grad_cam_plusplus import GradCAMPlusPlus
from pytorch_grad_cam.eigen_cam import EigenCAM
from pytorch_grad_cam.score_cam import ScoreCAM
from pytorch_grad_cam.xgrad_cam import XGradCAM
from pytorch_grad_cam import GradCAM
```

2. Defining necessary functions

In this section all fucntions which are necessary for this notebook are defined which include:

- image_transform()
- predict_and_get_target_class()
- get_conv_layer()

```
In [95]: def image_transform(image_url):
        """
        Transforms an image from a given URL into a preprocessed image tensor suitable for model inference.

        Parameters:
        - image_url (str): The URL of the image to be transformed.

        Returns:
        - img (numpy.ndarray): The resized image as a numpy array, normalized to the range [0, 1].
        - input_tensor (torch.Tensor): The preprocessed image tensor ready for model inference.

        Process:
        1. The image is loaded from the provided URL.
        2. The image is resized to 224x224 pixels.
        3. The pixel values are scaled to the range [0, 1].
        4. The image is normalized using the specified mean and standard deviation.
        """
        img = np.array(Image.open(requests.get(image_url, stream=True).raw))
        img = cv2.resize(img, (224, 224))
        img = np.float32(img) / 255
        input_tensor = preprocess_image(img, mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
        return img, input_tensor
```

```
In [96]: def predict_and_get_target_class(model, input_tensor):
        """
        Predicts the class of the given input tensor using the provided model and returns the predicted class index and label.

        Parameters:
        - model (torch.nn.Module): The neural network model used for making predictions.
        - input_tensor (torch.Tensor): The preprocessed image tensor from image_transform.

        Returns:
        - predicted_class (list): A list containing the predicted class index.
        - predicted_label (list): A list containing the predicted label name.

        Process:
        1. Passes the input tensor through the model to get the output.
        2. Determines the index of the highest value in the output tensor (using argmax).
        3. Returns the predicted class index as a list (list of one value).
        4. Maps the predicted class index to the corresponding human-readable label from the ImageNet labels.
        """
        # Fetch the ImageNet Labels
        response = requests.get('https://raw.githubusercontent.com/pytorch/hub/master/imagenet_classes.txt')
        labels = response.text.strip().split('\n')

        # Ensure the model is in evaluation mode
        model.eval()

        with torch.no_grad():
            # Get the model's output for the input tensor
            output = model(input_tensor)

            # Get the index of the highest value in the output tensor
            predicted_class = [torch.argmax(output).item()]

            # Get the predicted label name
            predicted_label = [labels[predicted_class[0]]]
```

```

        # Get the human-readable label name using the predicted class index
        predicted_label = labels[predicted_class[0]]

    return predicted_class, predicted_label

```

```

In [97]: def get_conv_layer(model):
        """
        Retrieves a specified convolutional layer from the given neural network model.

        Parameters:
        - model (torch.nn.Module): The neural network model to search through.

        Returns:
        - torch.nn.Conv2d: The selected convolutional layer based on user input.

        Process:
        1. Collects all convolutional layers (`torch.nn.Conv2d`) from the model.
        2. Prints the total number of convolutional layers found.
        3. Prompts the user to input the index of the convolutional layer they wish
        4. Checks if the user input is a valid layer number.
        5. Returns the convolutional layer.
        6. If the input is not valid, prints an error message.

        Note:
        - If there are no convolutional layers in the model, the function will not print anything.
        """
        conv_layers = [layer for layer in model.modules() if isinstance(layer, torch.nn.Conv2d)]
        input_layer = int(input(f'There are {len(conv_layers)} conv2d layers in this model. Enter the index of the layer you want: '))
        if input_layer <= 0 :
            return f'Please enter positive layer number'

        elif input_layer > len(conv_layers):
            return f'Please enter proper layer number'
        else:
            return conv_layers[input_layer - 1]

```

3. Importing image and model

In this section we use the above functions and also load the model and input image

```

In [98]: # storing the model in a variable.
        model = models.densenet121(pretrained=True)
        #model.eval()

```

```

In [99]: # storing the URL of the image in a variable.
        image_url='https://media.istockphoto.com/id/492696926/photo/wood-mouse-on-root-of-tree'
        img,input_tensor=image_transform(image_url)

```

3.1.1 Making prediction with the model

```

In [100]: # making prediction on the input tensor based on our model.
        index,name=predict_and_get_target_class(model,input_tensor)
        index,name

```

```
Out[100...] ([331], 'hare')
```

3.1.2 Saving the target label number

```
In [101...] # saving the prediction on the variable targets.(optional for Grad-CAM)
targets =[ClassifierOutputTarget(index)]
```

3.1.3 Getting the convolutional layer

```
In [102...] # Getting the target layer of the model for visualization.
target_layers = get_conv_layer(model)
target_layers
```

```
Out[102...] Conv2d(128, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
```

4. Applying the Gradcam methods

In this section all the CAM techniques are applied on model and input tensor which include:

- Grad-CAM
- Grad-CAM++
- Eigen-CAM
- XGrad-CAM

Note: These techniques are taken from jacob gill's github repository reference link:
<https://github.com/jacobgil/pytorch-grad-cam>

4.1 Normal Grad-CAM

```
In [103...] with GradCAM(model=model, target_layers=[target_layers]) as cam:
    grayscale_cams = cam(input_tensor=input_tensor, targets=targets)
    grad_cam_image = show_cam_on_image(img, grayscale_cams[0, :], use_rgb=True)
```

4.2 GradCAM++

```
In [104...] cam2=GradCAMPlusPlus(model=model, target_layers=[target_layers])
grayscale_cams2=cam2(input_tensor=input_tensor, targets=targets)
grayscale_cam2 = grayscale_cams2[0, :]
gradcamplusplus_image=show_cam_on_image(img,grayscale_cam2,use_rgb=True)
```

4.3 EigenCAM

```
In [105...] cam3=EigenCAM(model=model,target_layers=[target_layers])
grayscale_cams3=cam3(input_tensor=input_tensor)
grayscale_cam3=grayscale_cams3[0, :]
eigencam_image=show_cam_on_image(img,grayscale_cam3,use_rgb=True)
```

4.4 ScoreCAM

Here we have not shown the implementation of ScoreCAM because it takes a lot of time and the output is also not included in the results:

```
In [106... '''cam4=ScoreCAM(model=model,target_layers=target_layers)
grayscale_cams4=cam4(input_tensor=input_tensor)
grayscale_cam4=grayscale_cams4[0, :]
cam4_image=show_cam_on_image(img,grayscale_cam4,use_rgb=True)'''

Out[106... 'cam4=ScoreCAM(model=model,target_layers=target_layers)\ngrayscale_cams4=cam4(i
nput_tensor=input_tensor)\ngrayscale_cam4=grayscale_cams4[0, :]\ncam4_image=sho
w_cam_on_image(img,grayscale_cam4,use_rgb=True)'
```


4.5 XGradCAM

```
In [107... cam5=XGradCAM(model=model,target_layers=[target_layers])
grayscale_cams5=cam5(input_tensor=input_tensor)
grayscale_cam5=grayscale_cams5[0, :]
xgradcam_image=show_cam_on_image(img,grayscale_cam5,use_rgb=True)
```

5. visualizing

In this section the visulization takes place

```
In [108... # here we plot the visualizations and aslo print the label name of the based o
cam = np.uint8(255*grayscale_cams[0, :])
cam = cv2.merge([cam, cam, cam])
labels = ['Original', 'Grad-CAM', 'Grad-CAM++', 'Eigen-CAM', 'XGrad-CAM']
imagess = np.hstack((np.uint8(255*img), grad_cam_image, gradcamplusplus_image,ei
fig, axes = plt.subplots(1, len(labels), figsize=(15, 5))
for ax, image, label in zip(axes, [img, grad_cam_image, gradcamplusplus_image, e
    ax.imshow(image)
    ax.set_title(label)
    ax.axis('off')
plt.tight_layout()
plt.show()
print(f'The above image has {name} in it.')
```



The above image has hare in it.

The above picture highlights the parts of the image that the model focused on to make that prediction.

The highlighted part varies slightly between different techniques because each method uses distinct approaches to emphasize the regions.